

Background

Fully Bayes approach

Session information

# Fully Bayes Normal Means

*Matthew Stephens**May 3, 2018***Last updated:** 2018-05-03**workflowr checks:** (Click a bullet for more information)

- ► **✓ R Markdown file:** up-to-date
- ► **✓ Environment:** empty
- ► **✓ Seed:** `set.seed(20180411)`
- ► **✓ Session information:** recorded
- ► **✓ Repository version:** 196d0e3  
(<https://github.com/stephens999/stat34800/tree/196d0e34e4c5f85715f675878331d6554b754cb0>)

► **Expand here to see past versions:**

## Background

In a previous homework you implemented Empirical Bayes (EB) shrinkage for the normal means problem with a normal prior. That is we have data  $X = (X_1, \dots, X_n)$ :

$$X_j | \theta_j, s_j \sim N(\theta_j, s_j^2)$$

and assume

$$\theta_j | \mu, \sigma \sim N(\mu, \sigma^2) \quad j = 1, \dots, n.$$

The EB approach involved two steps:

1. Estimates  $\mu, \sigma$  by maximizing the log-likelihood  $l(\mu, \sigma) = \log p(X | \mu, \sigma)$ .
2. Compute the **posteriore distribution**  $p(\theta_j | \hat{\mu}, \hat{\sigma})$ .

The EB approach can be criticized for ignoring uncertainty in the estimates of  $\mu$  and  $\sigma$ . Here we will use MCMC to do a fully Bayesian analysis that takes account of this uncertainty.

## Fully Bayes approach

To make this easier we will first re-parameterize to use  $\eta = \log(\sigma)$ , so  $\eta$  can take any value on the real line.

We will use a uniform prior on  $(\mu, \eta)$ ,  $p(\mu, \eta) \propto 1$  in the range  $\mu \in [-a, a]$  and  $\eta \in [-b, b]$ . You can use  $a = 10^6$  and  $b = 10$ . (Because  $\eta$  is on the log scale,  $b = 10$  covers a wide range of possible standard deviations). Thus the posterior distribution on  $\mu, \eta$  is given by

$$p(\mu, \eta | X) \propto p(X | \mu, \eta) I(|\mu| < a) I(|\eta| < b)$$

where  $I$  denotes an indicator function.

1. Modify your **log-likelihood computation code** from your previous homework to compute the **log-likelihood for  $(\mu, \eta)$**  given data  $X$  (and standard deviations  $s$ ).
2. Use this to implement a **MH algorithm** to sample from  $\pi(\mu, \eta) \propto p(X | \mu, \eta)$ . Note: In computing the MH acceptance probability you need to compute a ratio  $L_1/L_2$ . For numerical stability reasons you should always compute this ratio by  $\exp(l_1 - l_2)$  where  $l_i = \log(L_i)$  rather than directly computing  $L_1$  and  $L_2$  and then computing their ratio. (If both  $L_1$  and  $L_2$  are very small, they may be 0 to machine precision, which causes problems if you try to compute  $L_1/L_2$  directly.)
3. Apply your **MH algorithm** to **simulated data** where you know the answer. Run your MH algorithm multiple (at least 3) times from multiple different initializations. For each run plot how the value of  $\log \pi(\mu^t, \eta^t)$  changes with iteration  $t$ . You should see that it starts from a low value (assuming you initialized to something that is not consistent with the data) and then gradually increases until it settles down to a “steady state” behavior. Use these plots to help decide **how many iterations to run** your algorithm to get reliable results (ie so results from different runs look similar) and **how many iterations to discard as “burn-in”**. Compare your posterior distributions of  $\mu$  and  $\eta$  with the true values you simulated (the distributions should cover the true values unless you did something wrong or are unlucky!).
4. Repeat part 3 for the “**8 schools data**” here (<http://andrewgelman.com/2014/01/21/everything-need-know-bayesian-statistics-learned-eight-schools/>) (omitting the comparisons with the true values, which of course you do not know here).
5. Note that the **posterior distribution on  $\theta_j$**  is given by:

$$p(\theta_j | X) = \int p(\theta_j | X, \mu, \eta) p(\mu, \eta | X)$$

which is the expectation of  $p(\theta_j | X, \mu, \eta)$  over the posterior  $p(\mu, \eta | X)$ . Computing **posterior distributions** like this is sometimes referred to as “**integrating out uncertainty in  $\mu, \eta$** ”. (It is useful to compare this with the EB approach of just plugging in the maximum likelihood estimates and computing  $p(\theta_j | X, \hat{\mu}, \hat{\eta})$ ). Notice that the **two will produce similar results if the posterior distribution  $p(\mu, \eta | X)$  is very concentrated around the mle.**)

Given  $T$  samples  $\mu^1, \eta^1, \dots, \mu^T, \eta^T$  from the posterior distribution  $p(\mu, \eta|X)$  you can approximate this expectation by

$$p(\theta_j|X) \approx (1/T) \sum_t p(\theta_j|X, \mu^t, \eta^t).$$

So you can approximate the **posterior mean** by

$$E(\theta_j|X) \approx (1/T) \sum_t E(\theta_j|X, \mu^t, \eta^t).$$

Using the same idea, given an **expression** to approximate the **posterior second moment**  $E(\theta_j^2|X)$ , and so approximate the **posterior variance** (and hence **posterior standard deviation**).

- Use the results from 4 and 5 to compute an approximate **posterior mean** and **posterior standard deviation** for  $\theta_j$  for each school in the **8 schools data**. Compare and contrast your results with the EB results and also the discussion in the initial blog-post here (<http://andrewgelman.com/2014/01/21/everything-need-know-bayesian-statistics-learned-eight-schools/>)

## Session information

```
sessionInfo()
```

```
R version 3.3.2 (2016-10-31)
Platform: x86_64-apple-darwin13.4.0 (64-bit)
Running under: OS X El Capitan 10.11.6

locale:
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

attached base packages:
[1] stats      graphics  grDevices  utils      datasets  methods   base

loaded via a namespace (and not attached):
[1] workflowr_1.0.1   Rcpp_0.12.16      digest_0.6.15
[4] rprojroot_1.3-2   R.methodsS3_1.7.1 backports_1.1.2
[7] git2r_0.21.0      magrittr_1.5       evaluate_0.10.1
[10] stringi_1.1.7     whisker_0.3-2     R.oo_1.22.0
[13] R.utils_2.6.0     rmarkdown_1.9     tools_3.3.2
[16] stringr_1.3.0     yaml_2.1.18       htmltools_0.3.6
[19] knitr_1.20
```

This reproducible R Markdown (<http://rmarkdown.rstudio.com>) analysis was created with workflowr (<https://github.com/jdblischak/workflowr>) 1.0.1

