

# #homework

- A: In this question you will extend the HMM in <https://stephens999.github.io/fiveMinuteStats/hmm.html> to treat the means of the two states as unknown and to be estimated. (Note that the true values of the means in the simulation are 1 and 2).
  - Derive and implement the **EM algorithm** for **estimating the means**.
  - Check your implementation by running it on the example and seeing that the log-likelihood is increasing. [Hint: note that the forwards algorithm gives you the likelihood.]
  - Try running the EM algorithm multiple times from different starting points. Does it get stuck in local optima of the log-likelihood?
- B: Complete the exercise in [https://stephens999.github.io/stat34800/hmm\\_exercise.html](https://stephens999.github.io/stat34800/hmm_exercise.html)
- C: **Spatial Gaussian Processes**
  - Consider the data from <http://journals.plos.org/plosbiology/article?id=10.1371/journal.pbio.0030339>, available at <https://github.com/stephens999/hgen48600/tree/master/data/CCR5>, which you can read into R using code from <https://stephens999.github.io/hgen48600/ccr5.html>
  - These consist of latitude, longitude, and an allele frequency at each location. We will model these data as a **Gaussian process**. Since allele frequency lies in  $[0,1]$  start by using the transformation  $x = \log(\hat{f}/(1 - \hat{f}))$ . (Here  $\hat{f}$  is the estimated frequency in the code above.) We will let  $y$  denote locations in space (latitude, longitude) and  $x(\cdot)$  denote the allele frequency varying as a function of space, so  $x(y) \in [0,1]$  is the allele frequency at location  $y$ . We will model  $x(\cdot)$  as a Gaussian process, with constant mean  $\mu = m$  and squared exponential covariance function of the form  $a_1 \exp(-(d/a_2)^2)$ .
 

$\hat{f} \in [0,1]$   
 $(-\infty, \infty)$
  - Hence,  $a = (a_1, a_2)$  and the mean  $m$  are the parameters to be estimated.
  - Write a function to compute the covariance matrix for  $x^{\text{obs}} := (x(y_1), \dots, x(y_r))$  given a value of  $a$ . Here  $y_1, \dots, y_r$  are the locations at which you have observations in the dataset. Try a few values of  $a$  and check that the resulting covariance matrix is valid - that is, it is positive semi-definite. (The best way to check that a covariance is positive semi-definite is to attempt to perform a cholesky decomposition: if the decomposition succeeds then the matrix must be PSD).
  - Write a function to compute the log-likelihood for the data  $x^{\text{obs}}$ , given  $a, m$ . [Here we assume the mean is constant across the whole region, so  $m$  is the same at every location].

- The model here is that  $(x^{\text{obs}} | m, a) \sim N_r(\mu, \Sigma)$  where  $\Sigma = \Sigma(a)$  is the function of  $a$  that you coded above and  $\mu = \text{rep}(m, r)$ . So your likelihood just involves computing a multivariate normal density. You can use the R function `mvtnorm::dmvnorm` (with `log=TRUE`)
- Try using the R function `optim` (or another approach if you prefer) to optimize the likelihood numerically over  $a, m$ . (I found it seemed to work OK, in that it gave similar answers from different starting points).
- Now we are going to try deleting each of the observed data points in turn and "impute" its value using our model. This process is sometimes known as **Kriging**.
  - Let  $X = (X_1, \dots, X_r)$  be  $r$ -variate normal with mean  $\mu$  and variance covariance  $\Sigma$ . Write a function to compute the conditional expectation of  $X_1$  given  $X_2, \dots, X_r$ . [This is an application of standard results for the **conditional mean of a Gaussian** from, e.g. [https://en.wikipedia.org/wiki/Multivariate\\_normal\\_distribution#Conditional\\_distributions](https://en.wikipedia.org/wiki/Multivariate_normal_distribution#Conditional_distributions)]
  - Apply this function to compute  $E(x(y_1) | x(y_2), \dots, x(y_r))$ . Notice that this expectation ends up being a weighted linear combination of the other datapoints. Intuitively, if allele frequencies **vary smoothly in space** then this weighted linear combination should **weight the nearby data points more**. Does it?
  - Repeat this for each of the  $r$  datapoints.
  - How does the accuracy of this imputation scheme compare with just using the **mean of the other datapoints** to impute each datapoint?